

# Boston house price prediction

## Importing libraries:

In [31]:

```
# Import

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import metrics
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
```

## Boston dataset from sklearn

In [3]:

```
from sklearn.datasets import load_boston

boston = load_boston()
```

In [4]:

```
data = pd.DataFrame(boston.data)

print(data)
```

	0	1	2	3	4	5	6	7	8	9	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	
5	0.02985	0.0	2.18	0.0	0.458	6.430	58.7	6.0622	3.0	222.0	
6	0.08829	12.5	7.87	0.0	0.524	6.012	66.6	5.5605	5.0	311.0	
7	0.14455	12.5	7.87	0.0	0.524	6.172	96.1	5.9505	5.0	311.0	
8	0.21124	12.5	7.87	0.0	0.524	5.631	100.0	6.0821	5.0	311.0	
9	0.17004	12.5	7.87	0.0	0.524	6.004	85.9	6.5921	5.0	311.0	
10	0.22489	12.5	7.87	0.0	0.524	6.377	94.3	6.3467	5.0	311.0	
11	0.11747	12.5	7.87	0.0	0.524	6.009	82.9	6.2267	5.0	311.0	
12	0.09378	12.5	7.87	0.0	0.524	5.889	39.0	5.4509	5.0	311.0	
13	0.62976	0.0	8.14	0.0	0.538	5.949	61.8	4.7075	4.0	307.0	
14	0.63796	0.0	8.14	0.0	0.538	6.096	84.5	4.4619	4.0	307.0	
15	0.62739	0.0	8.14	0.0	0.538	5.834	56.5	4.4986	4.0	307.0	
16	1.05393	0.0	8.14	0.0	0.538	5.935	29.3	4.4986	4.0	307.0	
17	0.78420	0.0	8.14	0.0	0.538	5.990	81.7	4.2579	4.0	307.0	
18	0.80271	0.0	8.14	0.0	0.538	5.456	36.6	3.7965	4.0	307.0	
19	0.72580	0.0	8.14	0.0	0.538	5.727	69.5	3.7965	4.0	307.0	
20	1.25179	0.0	8.14	0.0	0.538	5.570	98.1	3.7979	4.0	307.0	
21	0.85204	0.0	8.14	0.0	0.538	5.965	89.2	4.0123	4.0	307.0	
22	1.23247	0.0	8.14	0.0	0.538	6.142	91.7	3.9769	4.0	307.0	
23	0.98843	0.0	8.14	0.0	0.538	5.813	100.0	4.0952	4.0	307.0	
24	0.75026	0.0	8.14	0.0	0.538	5.924	94.1	4.3996	4.0	307.0	
25	0.84054	0.0	8.14	0.0	0.538	5.599	85.7	4.4546	4.0	307.0	
26	0.67191	0.0	8.14	0.0	0.538	5.813	90.3	4.6820	4.0	307.0	
27	0.95577	0.0	8.14	0.0	0.538	6.047	88.8	4.4534	4.0	307.0	
28	0.77299	0.0	8.14	0.0	0.538	6.495	94.4	4.4547	4.0	307.0	
29	1.00245	0.0	8.14	0.0	0.538	6.674	87.3	4.2390	4.0	307.0	
...	...	...	...	...	...	...	...	...	...	...	
476	4.87141	0.0	18.10	0.0	0.614	6.484	93.6	2.3053	24.0	666.0	

477	15.02340	0.0	18.10	0.0	0.614	5.304	97.3	2.1007	24.0	666.0
478	10.23300	0.0	18.10	0.0	0.614	6.185	96.7	2.1705	24.0	666.0
479	14.33370	0.0	18.10	0.0	0.614	6.229	88.0	1.9512	24.0	666.0
480	5.82401	0.0	18.10	0.0	0.532	6.242	64.7	3.4242	24.0	666.0
481	5.70818	0.0	18.10	0.0	0.532	6.750	74.9	3.3317	24.0	666.0
482	5.73116	0.0	18.10	0.0	0.532	7.061	77.0	3.4106	24.0	666.0
483	2.81838	0.0	18.10	0.0	0.532	5.762	40.3	4.0983	24.0	666.0
484	2.37857	0.0	18.10	0.0	0.583	5.871	41.9	3.7240	24.0	666.0
485	3.67367	0.0	18.10	0.0	0.583	6.312	51.9	3.9917	24.0	666.0
486	5.69175	0.0	18.10	0.0	0.583	6.114	79.8	3.5459	24.0	666.0
487	4.83567	0.0	18.10	0.0	0.583	5.905	53.2	3.1523	24.0	666.0
488	0.15086	0.0	27.74	0.0	0.609	5.454	92.7	1.8209	4.0	711.0
489	0.18337	0.0	27.74	0.0	0.609	5.414	98.3	1.7554	4.0	711.0
490	0.20746	0.0	27.74	0.0	0.609	5.093	98.0	1.8226	4.0	711.0
491	0.10574	0.0	27.74	0.0	0.609	5.983	98.8	1.8681	4.0	711.0
492	0.11132	0.0	27.74	0.0	0.609	5.983	83.5	2.1099	4.0	711.0
493	0.17331	0.0	9.69	0.0	0.585	5.707	54.0	2.3817	6.0	391.0
494	0.27957	0.0	9.69	0.0	0.585	5.926	42.6	2.3817	6.0	391.0
495	0.17899	0.0	9.69	0.0	0.585	5.670	28.8	2.7986	6.0	391.0
496	0.28960	0.0	9.69	0.0	0.585	5.390	72.9	2.7986	6.0	391.0
497	0.26838	0.0	9.69	0.0	0.585	5.794	70.6	2.8927	6.0	391.0
498	0.23912	0.0	9.69	0.0	0.585	6.019	65.3	2.4091	6.0	391.0
499	0.17783	0.0	9.69	0.0	0.585	5.569	73.5	2.3999	6.0	391.0
500	0.22438	0.0	9.69	0.0	0.585	6.027	79.7	2.4982	6.0	391.0
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0

	10	11	12
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33
5	18.7	394.12	5.21
6	15.2	395.60	12.43
7	15.2	396.90	19.15
8	15.2	386.63	29.93
9	15.2	386.71	17.10
10	15.2	392.52	20.45
11	15.2	396.90	13.27
12	15.2	390.50	15.71
13	21.0	396.90	8.26
14	21.0	380.02	10.26
15	21.0	395.62	8.47
16	21.0	386.85	6.58
17	21.0	386.75	14.67
18	21.0	288.99	11.69
19	21.0	390.95	11.28
20	21.0	376.57	21.02
21	21.0	392.53	13.83
22	21.0	396.90	18.72
23	21.0	394.54	19.88
24	21.0	394.33	16.30
25	21.0	303.42	16.51
26	21.0	376.88	14.81
27	21.0	306.38	17.28
28	21.0	387.94	12.80
29	21.0	380.23	11.98
..	...	...	...
476	20.2	396.21	18.68
477	20.2	349.48	24.91
478	20.2	379.70	18.03
479	20.2	383.32	13.11
480	20.2	396.90	10.74
481	20.2	393.07	7.74
482	20.2	395.28	7.01
483	20.2	392.92	10.42
484	20.2	370.73	13.34
485	20.2	388.62	10.58

```
486 20.2 392.68 14.98
487 20.2 388.22 11.45
488 20.1 395.09 18.06
489 20.1 344.05 23.97
490 20.1 318.43 29.68
491 20.1 390.11 18.07
492 20.1 396.90 13.35
493 19.2 396.90 12.01
494 19.2 396.90 13.59
495 19.2 393.29 17.60
496 19.2 396.90 21.14
497 19.2 396.90 14.10
498 19.2 396.90 12.92
499 19.2 395.77 15.10
500 19.2 396.90 14.33
501 21.0 391.99 9.67
502 21.0 396.90 9.08
503 21.0 396.90 5.64
504 21.0 393.45 6.48
505 21.0 396.90 7.88
```

```
[506 rows x 13 columns]
```

## Notations (Those notations are copied from the source):

**1- CRIM** per capita crime rate by town **2- ZN** proportion of residential land zoned for lots over 25,000 sq.ft. **3- INDUS** proportion of non-retail business acres per town **4- CHAS** Charles River dummy variable (= 1 if tract bounds river; 0 otherwise) **5- NOX** nitric oxides concentration (parts per 10 million) **6- RM** average number of rooms per dwelling **7- AGE** proportion of owner-occupied units built prior to 1940 **8- DIS** weighted distances to five Boston employment centres **9- RAD** index of accessibility to radial highways **10- TAX** full-value property-tax rate per 10,000usd **11- PTRATIO** pupil-teacher ratio by town **12- B**  $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of blacks by town **13- LSTAT** % lower status of the population

```
In [5]:
```

```
# now, we need to add what each column corresponds to:
```

```
data.columns = boston.feature_names
data.head()
```

```
Out[5]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [6]:
```

```
# specifing our target
```

```
data["Price"] = boston.target
```

```
In [7]:
```

```
# Properties:
```

```
print(data.shape)
```

```
print(data.columns)
```

```
print(data.dtypes)
```

```
print(data.nunique())
```

```
print(data.isnull().sum())
```

```
print(data.describe())
```

```
(506, 14)
```

```
Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',  
      'PTRATIO', 'B', 'LSTAT', 'Price'],  
      dtype='object')
```

```
CRIM      float64  
ZN        float64  
INDUS     float64  
CHAS      float64  
NOX       float64  
RM        float64  
AGE       float64  
DIS       float64  
RAD       float64  
TAX       float64  
PTRATIO   float64  
B         float64  
LSTAT     float64  
Price     float64
```

```
dtype: object
```

```
CRIM      504  
ZN        26  
INDUS     76  
CHAS       2  
NOX       81  
RM       446  
AGE      356  
DIS      412  
RAD        9  
TAX       66  
PTRATIO   46  
B       357  
LSTAT   455  
Price   229
```

```
dtype: int64
```

```
CRIM      0  
ZN        0  
INDUS     0  
CHAS      0  
NOX       0  
RM        0  
AGE       0  
DIS       0  
RAD       0  
TAX       0  
PTRATIO   0  
B         0  
LSTAT     0  
Price     0
```

```
dtype: int64
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	3.593761	11.363636	11.136779	0.069170	0.554695	6.284634	
std	8.596783	23.322453	6.860353	0.253994	0.115878	0.702617	
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	
75%	3.647423	12.500000	18.100000	0.000000	0.624000	6.623500	
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	

	AGE	DIS	RAD	TAX	PTRATIO	B	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032	
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864	
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000	
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500	

50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	LSTAT	Price
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

In [8]:

```
# Creating correlation between the features
```

```
correlation = data.corr()
```

```
correlation.shape
```

Out[8]:

```
(14, 14)
```

In [9]:

```
# Plotting heat map of the correlated features:
```

```
plt.figure(figsize = (20,20))
```


```
print("\n This is the Heatmap: \n")
```


```
sns.heatmap(correlation, cbar=True, square=True, fmt='.1f', annot=True, annot_kws={'size':15}, cmap='Greens')
```

This is the Heatmap:

Out[9]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x14e1852d5f8>
```





In [10]:

```
# specifying target variable and independent variable:
```

```
X = data.drop(['Price'], axis = 1)
y = data['Price']
```

In [11]:

```
# specifying training and testing data
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 4)
```

**Now, we are going to create the ML path, using Linear Regression and Random Forest Model:**

## Linear Regression:

### Training data:

In [12]:

```
# creating our model
```

```
model = LinearRegression()
```

```
# using the model on the training data
```

```
model.fit(X_train, y_train)
```

Out[12]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

In [14]:

```
# the intercept
```

```
model.intercept_
```

Out[14]:

```
36.357041376595035
```

In [18]:

```
# coefficients:
```

```
coeff = pd.DataFrame([X_train.columns,model.coef_]).T
coeff = coeff.rename(columns = {0:"Attributes",1:"Coefficients"})

coeff
```

Out[18]:

	Attributes	Coefficients
0	CRIM	-0.12257
1	ZN	0.0556777
2	INDUS	-0.00883428
3	CHAS	4.69345
4	NOX	-14.4358
5	RM	3.28008
6	AGE	-0.00344778
7	DIS	-1.55214
8	RAD	0.32625
9	TAX	-0.0140666
10	PTRATIO	-0.803275
11	B	0.00935369
12	LSTAT	-0.523478

In [24]:

```
# Evaluating the Model

y_pred = model.predict(X_train)

pd.DataFrame(y_pred)
```

Out[24]:

	0
0	24.522480
1	15.197510
2	25.577206
3	13.939400
4	39.466513
5	17.459599
6	39.710299
7	16.517481
8	20.197333
9	40.797755
10	33.572450
11	14.504206
12	11.445145
13	23.065640
14	24.397344
15	25.010961
16	14.361165
17	28.283415

18	25.049319
19	22.428252
20	21.815885
21	18.852087
22	13.356212
23	13.657927
24	23.647660
25	18.068763
26	16.129572
27	41.124149
28	19.433918
29	13.179809
...	...
324	15.155495
325	17.994417
326	30.761722
327	29.543075
328	6.259368
329	27.179352
330	14.805956
331	23.594438
332	22.668606
333	16.020712
334	24.058107
335	20.661335
336	25.379353
337	27.553691
338	26.950710
339	26.755668
340	19.869935
341	19.690256
342	24.332599
343	21.924869
344	20.354469
345	35.338450
346	13.007641
347	25.813350
348	22.959968
349	8.608369
350	31.511078
351	13.647191
352	26.501062
353	20.540965

354 rows × 1 columns



## Metrics:

In [25]:


```
print('R^2:', metrics.r2_score(y_train, y_pred))
print('Adjusted R^2:', 1 - (1 - metrics.r2_score(y_train, y_pred)) * (len(y_train) - 1) / (len(y_train) - X_train.shape[1] - 1))
print('MAE:', metrics.mean_absolute_error(y_train, y_pred))
print('MSE:', metrics.mean_squared_error(y_train, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

```
R^2: 0.7465991966746854
Adjusted R^2: 0.736910342429894
MAE: 3.0898610949711287
MSE: 19.073688703469028
RMSE: 4.367343437774161
```

## Plotting:

In [26]:

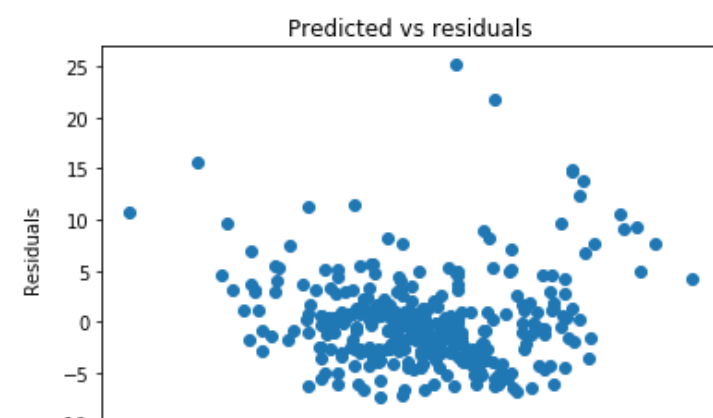
```
plt.scatter(y_train, y_pred)
plt.xlabel("Prices")
plt.ylabel("Predicated Prices")
plt.title("Predicted Prices VS Prices")
plt.show()
```




## Testing for residuals:

In [27]:

```
plt.scatter(y_pred, y_train - y_pred)
plt.title("Predicted vs residuals")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.show()
```





## Now, we need to use test data to validate our LR model:

In [29]:

```
y_test_pred = model.predict(X_test)
```

In [30]:

```
# Evaluating our model

acc_linreg = metrics.r2_score(y_test, y_test_pred)
print('R^2:', acc_linreg)
print('Adjusted R^2:', 1 - (1-metrics.r2_score(y_test, y_test_pred))*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1))
print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

```
R^2: 0.7120461624218655
Adjusted R^2: 0.6849200762732006
MAE: 3.867069394655806
MSE: 30.06816053374662
RMSE: 5.483444221814117
```

- The results are similar to the of the train data. As such, we are not over fitting or under fitting our model.

## Random Forest Model (RFM)

In [32]:

```
# initiating our RFM

rfm = RandomForestRegressor()

# Training our data

rfm.fit(X_train, y_train)
```

Out[32]:

```
RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                        max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
                        oob_score=False, random_state=None, verbose=0, warm_start=False)
```

In [33]:

```
# Evaluating our model

y_pred = rfm.predict(X_train)
```

In [34]:

```
# Model Evaluation
print('R^2:', metrics.r2_score(y_train, y_pred))
print('Adjusted R^2:', 1 - (1-metrics.r2_score(y_train, y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
print('MAE:', metrics.mean_absolute_error(y_train, y_pred))
print('MSE:', metrics.mean_squared_error(y_train, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

```
R^2: 0.9729121848992065
```

Adjusted R<sup>2</sup>: 0.9718764743218232  
MAE: 0.8334180790960453  
MSE: 2.038922316384181  
RMSE: 1.4279083711443745

## Plotting:


In [35]:

```
plt.scatter(y_train, y_pred)
plt.xlabel("Prices")
plt.ylabel("Predicted prices")
plt.title("Prices vs Predicted prices")
plt.show()
```



In [36]:

```
plt.scatter(y_pred, y_train-y_pred)
plt.title("Predicted vs residuals")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.show()
```



## Testing our the validity of our model:

In [37]:

```
y_pred_test = rfm.predict(X_test)
```

In [38]:

```
# Model Evaluation
```

```
acc_rf = metrics.r2_score(y_test, y_test_pred)
print('R^2:', acc_rf)
print('Adjusted R^2:', 1 - (1-metrics.r2_score(y_test, y_test_pred))*(len(y_test)-1)/(len
(y_test)-X_test.shape[1]-1))
print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

```
R^2: 0.7120461624218655
Adjusted R^2: 0.6849200762732006
MAE: 3.867069394655806
MSE: 30.06816053374662
RMSE: 5.483444221814117
```